

# Introduction to Machine Learning

M390

Dave Miller

1/23/2019

# Machine Learning

- **Machine learning** is a data analytics technique that teaches computers to do what comes naturally to humans and animals: learn from experience.
- Algorithms use computational methods to “learn” information directly from data without relying on a predetermined equation as a model.

# Types of Machine Learning

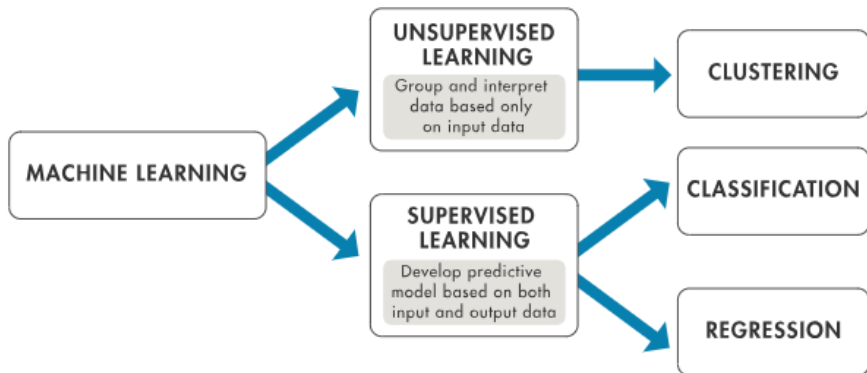


Figure 1. Machine learning techniques include both unsupervised and supervised learning.

**Figure 1:**

# Supervised Learning

- **Supervised machine learning** builds a model that makes predictions based on evidence in the presence of uncertainty.
- A supervised learning algorithm takes a known set of indicators (input data) and **known responses** to the data (output) and trains a model to generate reasonable predictions for the response to new data:

$$Y \sim f(X)$$

- Use supervised learning if you have known data for the output you are trying to predict.

# Regression

- **Regression** techniques predict continuous responses:
  - ▶ House (or anything else) price,
  - ▶ Scores of a game,
  - ▶ Test score,
  - ▶ etc.
- Use regression techniques if you are working with a data range or if the nature of your response is a real number.
- Common regression algorithms include linear regression, nonlinear models, regularization, stepwise regression, boosted and bagged decision trees, neural networks.

# Classification

- Classification techniques predict discrete responses:
  - ▶ whether a tumor is cancerous or benign,
  - ▶ Yes or No,
  - ▶ Win or loss,
  - ▶ Pass or fail.
  - ▶ etc.
- Use classification if your data can be tagged, categorized, or separated into specific groups or classes.
- Common classification algorithms include logistic regression, support vector machine (SVM), boosted and bagged decision trees and neural networks.

# Unsupervised Learning: Clustering

- **Unsupervised learning** finds hidden patterns or intrinsic structures in data. It is used to draw inferences from datasets consisting of input data **without labeled responses**.
- **Clustering** is the most common unsupervised learning technique. It is used for exploratory data analysis to find hidden patterns or groupings in data.



# Preparing for Regression Problems: R packages

- We will use the tidyverse package extensively throughout the course. More information can be found at <http://www.tidyverse.org/>.
- To install tidyverse, use `install.packages(tidyverse)`. If you already have tidyverse installed, run `library(tidyverse)`. The first two lines of all of my R Scripts look like the following:

```
#install.packages(tidyverse)  
library(tidyverse)
```



# Preparing for Regression Problems: Getting data

- Preloaded data in R:
  - ▶ View all by typing `data()`
  - ▶ Example: `mtcars`
- Loading your own data from my GitHub:

```
ex_1 <-  
read.csv('https://raw.githubusercontent.com/davmiller/M390/master/  
data/example_data_1.csv')
```

# Preparing for Regression Problems: Looking at the data

```
# Look at the first 6 rows.
```

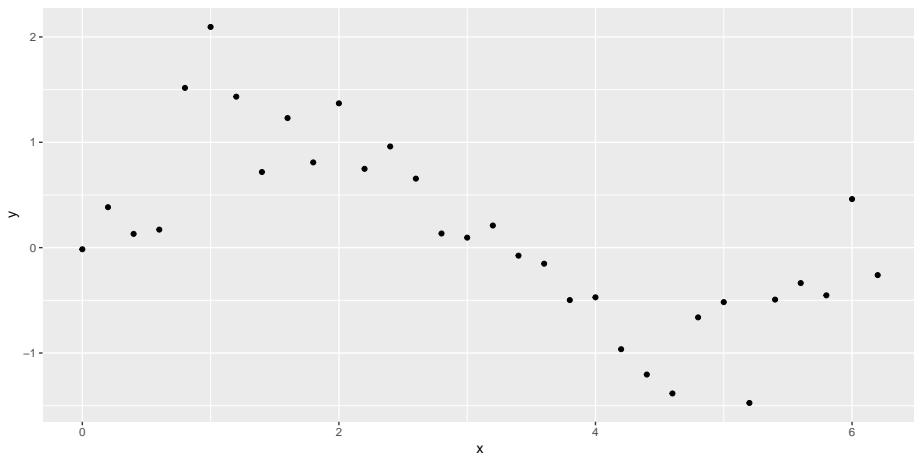
```
head(ex_1)
```

```
##           x           y
## 1 0.0 -0.01514382
## 2 0.2  0.38408346
## 3 0.4  0.13086018
## 4 0.6  0.17121048
## 5 0.8  1.51642042
## 6 1.0  2.09502139
```

- View(ex\_1) to view entire dataset.

# Plotting the data

```
# Plot using ggplot  
ggplot(data=ex_1, aes(x=x,y=y))+  
  geom_point()
```



# Preparing for Regression Problems: Data Splitting

- Main goal of regression is to find an algorithm  $f(X)$  that most accurately predicts future values (or responses)  $Y$  based on a set of inputs (or indicators)  $X$ :

$$Y \sim f(X)$$

- We want an algorithm that not only fits well to our past data, but more importantly, one that predicts a future outcome accurately. This is called the **generalizability** of our algorithm.

# Data Splitting

- To provide an accurate understanding of the generalizability of our final optimal model, we split our data into training and test data sets:
  - ▶ **Training Set:** these data are used to train our algorithms and tune hyper-parameters.
  - ▶ **Test Set:** having chosen a final model, these data are used to estimate its prediction error (generalization error). These data should not be used during model training!
- Given a fixed amount of data, typical recommendations for splitting your data into training-testing splits include 60% (training) - 40% (testing), 70%-30%, or 80%-20%.

# Data Splitting



**Figure 3:** Splitting data into training and test sets

# Simple random splitting

- The simplest way to split the data into training and test sets is to take a simple random sample.

```
# 70% random train / test split  
set.seed(123)  
index    <- sample(1:nrow(ex_1), round(nrow(ex_1) * 0.7))  
train_1  <- ex_1[index, ]  
test_1   <- ex_1[-index, ]
```

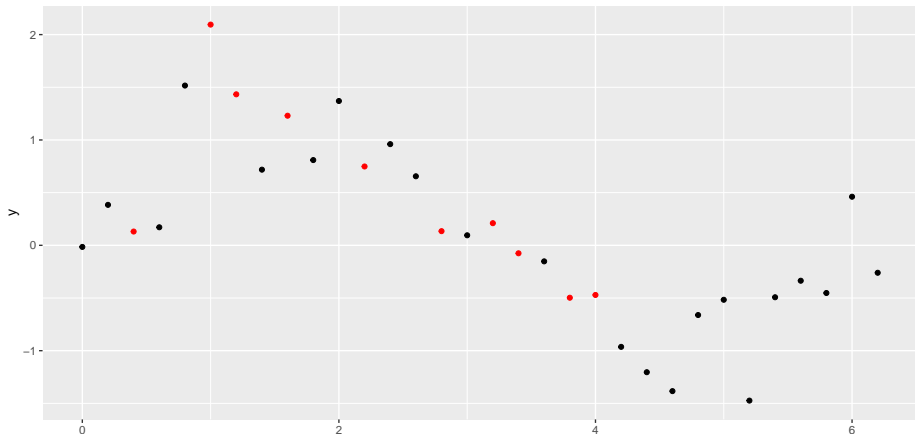
# Plotting train / test sets

*# Train in black, test in red.*

```
ggplot()+
```

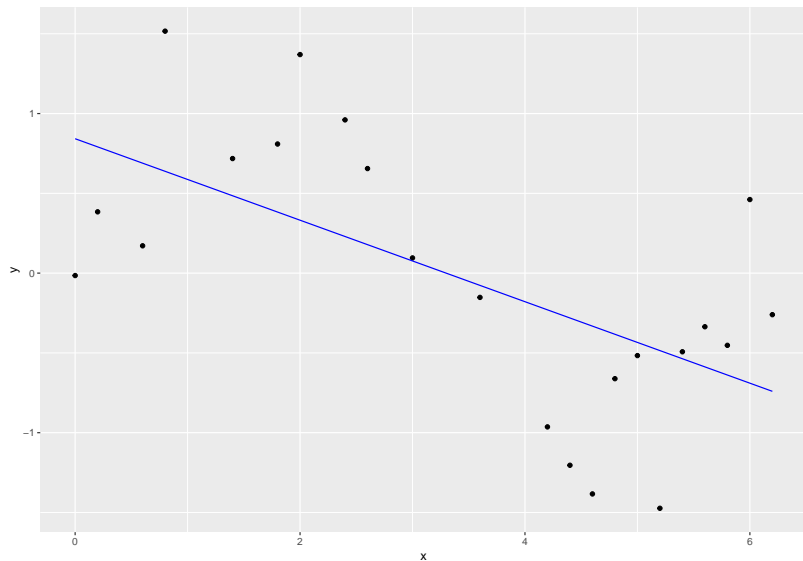
```
  geom_point(data=train_1, aes(x=x,y=y), color='black')+ 
```

```
  geom_point(data=test_1, aes(x=x,y=y), color='red')
```

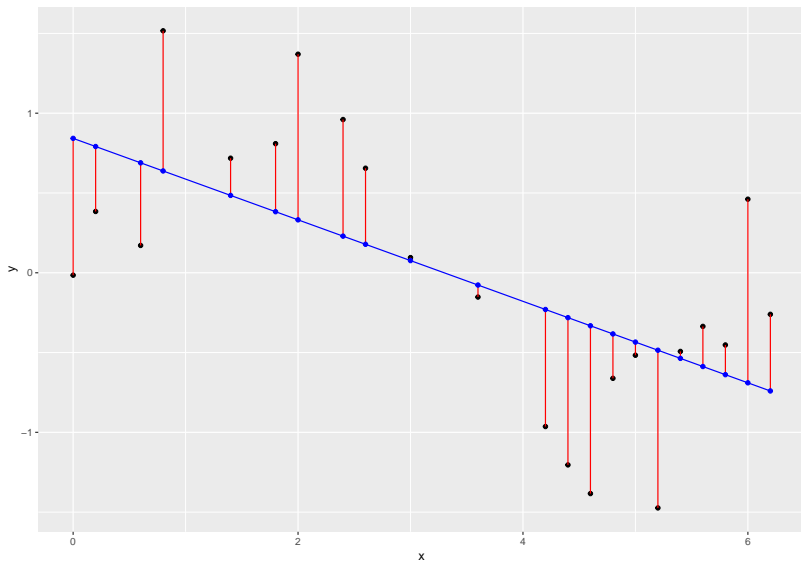




# Models trained on training set



# Evaluating a regression model



```
## [1] "Train MSE: 0.4197"
```

# Evaluating a regression model

- Mean squared error (**MSE**) is the average of the squared error.

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2.$$

**Objective: minimize**

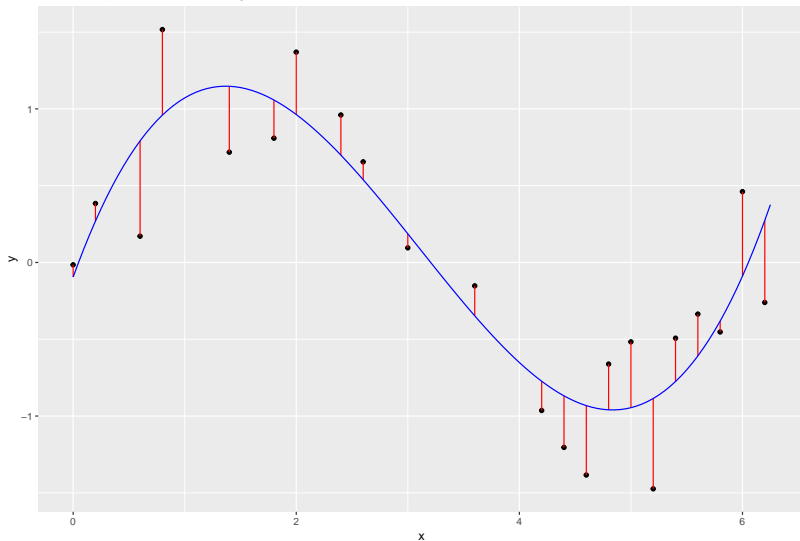
- Root mean squared error (**RMSE**) simply takes the square root of the MSE:

$$RMSE = \sqrt{MSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - (\beta_0 + \beta_1 x_i))^2}.$$

**Objective: minimize**

# Other models

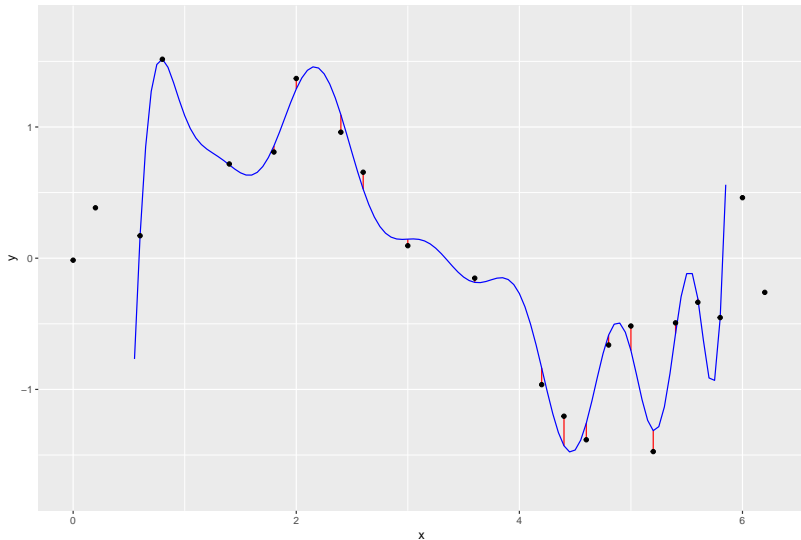
More complex model on training set



```
## [1] "Train MSE: 0.1351"
```

# Other models

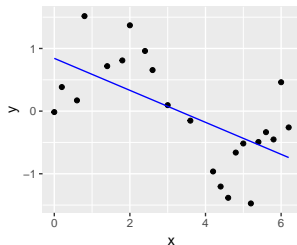
Even more complex model on training set



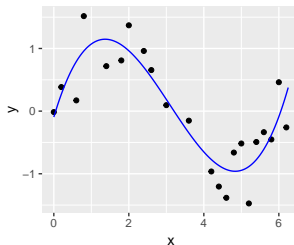
```
## [1] "Train MSE: 0.0092"
```

# Tuning a model

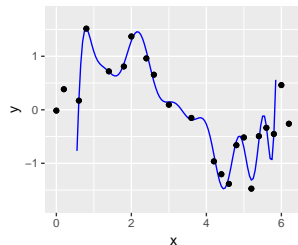
Too restrictive



Flexible



Overfit



# Overfitting

- **Overfitting** refers to a model that models the training data too well.
- Overfitting occurs when the evaluation metric on the training set is much better than on the test set.
- Underfitting refers to a model that can neither model the training data nor generalize to new data.

# Overfitting

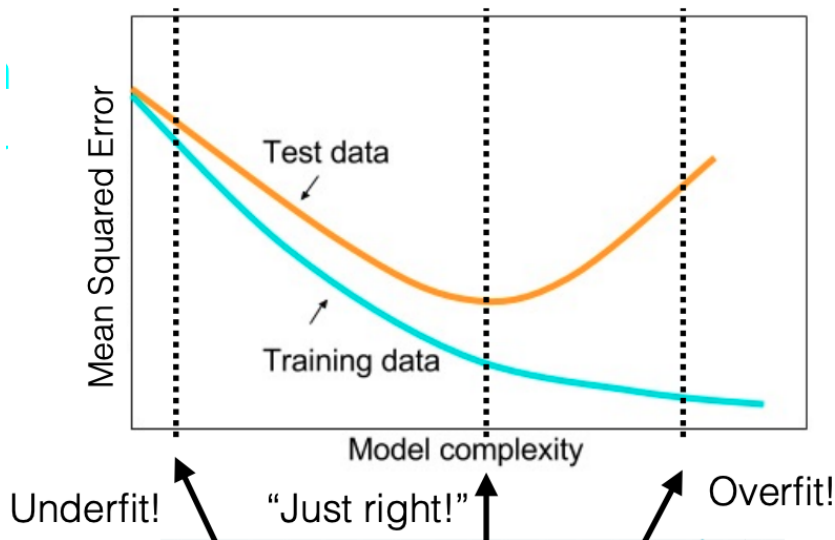
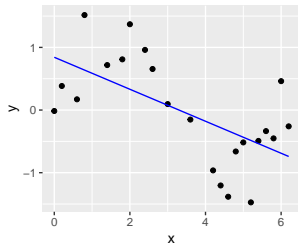


Figure 4: Finding the balance

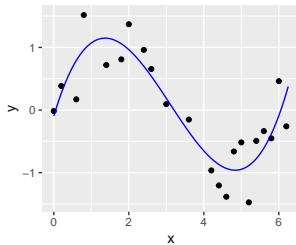


# Overfitting

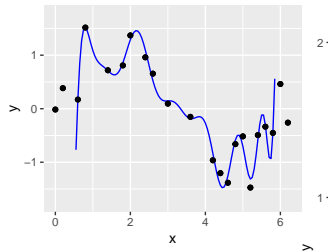
Train MSE: 0.4197



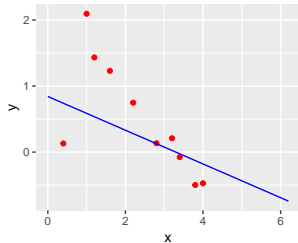
Train MSE: 0.1351



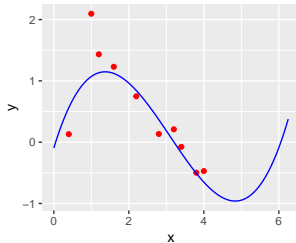
Train MSE: 0.0092



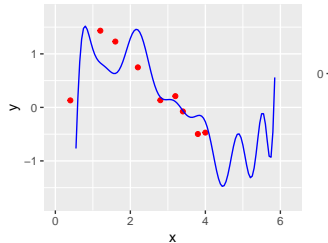
Test MSE: 0.4562



Test MSE: 0.1484



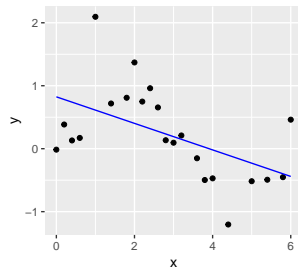
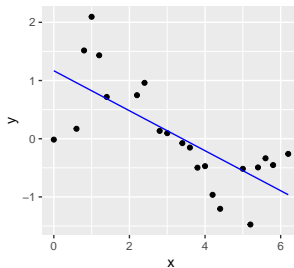
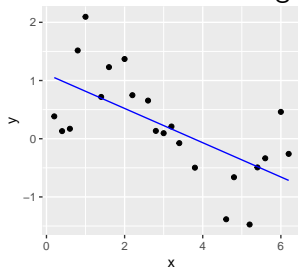
Test MSE: 2.0706



# Bias-variance tradeoff

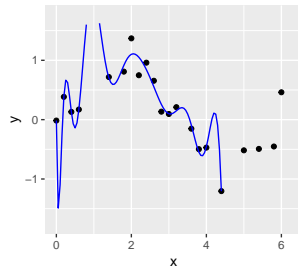
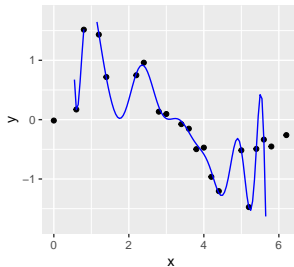
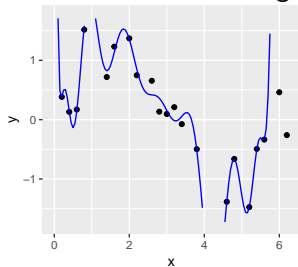
- A model that is rigid and consistent, which doesn't change much at all when provided a new training sample is said to have high **bias**.
- A model that changes significantly with small changes to the training sample is said to have high **variance**.

Three different training sets from same base set.



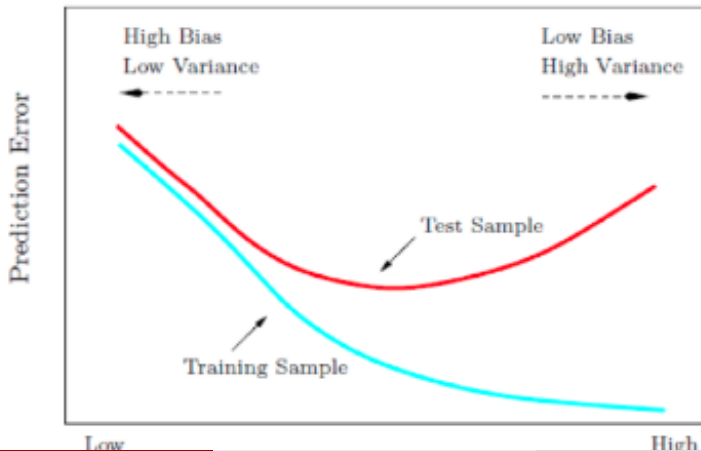
# Variance

Three different training sets from same base set.



# Bias-variance tradeoff

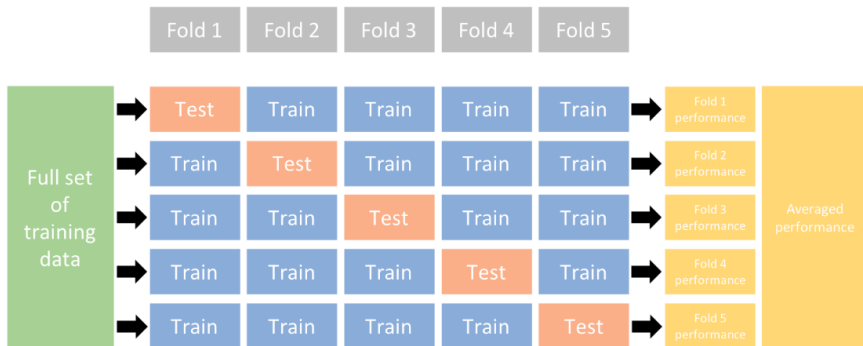
We want a model that balances bias and variance, and likely will minimize the error on future unseen data compared to the high bias or high variance models.



# Cross validation

- To find the model that balances the bias-variance tradeoff, we search for a model that minimizes a  **$k$ -fold cross-validation error metric**.
- $k$ -fold cross-validation is a resampling method that randomly divides the training data into  $k$  groups (aka folds) of approximately equal size.
- The model is fit on  $k - 1$  folds and then the held-out validation fold is used to compute the error.
- This procedure is repeated  $k$  times; each time, a different group of observations is treated as the validation set.
- This process results in  $k$  estimates of the test error, which are averaged to compute the  $k$ -fold CV estimate, which provides us with an approximation of the error to expect on unseen data.

# Cross validation



**Figure 6:** Cross validation